


Lecture 9

CMPEN 331

A decorative blue graphic consisting of a curved line that starts near the top left and sweeps downwards and to the right, forming a large, solid blue area in the bottom right corner of the slide.


Jump Addressing

- Jump (j and jal) targets could be anywhere in text segment
 - Encode full address in instruction



Branch Addressing

- Branch instructions specify
 - Opcode, two registers, target address
- Most branch targets are near branch
 - Forward or backward



- PC (program counter)-relative addressing
 - Target address = PC + branch address

Target Addressing Example

- Loop code from earlier example
 - Assume Loop at location 80000

```

Loop: sll    $t1, $s3, 2      80000
      add    $t1, $t1, $s6    80004
      lw     $t0, 0($t1)      80008
      bne    $t0, $s5, Exit    80012
      addi   $s3, $s3, 1      80016
      j      Loop             80020
Exit: ...                     80024
    
```

0	0	19	9	2	0
0	9	22	9	0	32
35	9	8	0		
5	8	21	2		
8	19	19	1		
2	20000				

Branching Far Away

- If branch target is too far to encode with 16-bit offset, assembler rewrites the code
- Example

```
        beq $s0,$s1, L1
          ↓
        bne $s0,$s1, L2
        j  L1
L2:     ...
```

Addressing Modes

- Addressing modes are the ways of specifying an operand or a memory address.
- It is how an address (memory or register) is determined.
- Instruction type is how the instruction is put together.
- Example: addi, beq, and lw are all I-types instructions.
 - addi uses immediate addressing mode
 - beq uses pc-relative addressing
 - lw uses base addressing

Addressing Modes

- **IMMEDIATE:** a numeric value embedded in the instruction is the actual operand.
- **REGISTER:** a source or destination operand is specified as content of one of the registers `$0-$31`. This is used in the **jr** (jump register) instruction
- **PC-RELATIVE:** a data or instruction memory location is specified as an offset relative to the incremented PC. This is used in the **beq** and **bne** (branch equal, branch not equal) instructions.
- **BASE:** a data or instruction memory location is specified as assigned offset from a register. This is used in the **lw** and **sw** (load word, store word) instructions.
- **PSEUDO-DIRECT:** the memory address is (mostly) embedded in the instruction. This is used in the **j** (jump) instruction.

Addressing Modes

- Register addressing

Operand is in register

add \$s1, \$s2, \$s3 means $\$s1 \leftarrow \$s2 + \$s3$

- Base addressing

Operand is in memory.

The address is the sum of a register and a constant.

lw \$s1, 32(\$s3) means $\$s1 \leftarrow M[s3 + 32]$

- Immediate addressing

The operand is a constant.

addi \$s1, \$zero, 7 means $\$s1 \leftarrow 0 + 7$

- PC-relative addressing

The operand address = PC + an offset

Implements position-independent codes.

- Pseudo-direct addressing

Used in the J format. The target address is the concatenation of the 4 MSB's of the PC with the 28-bit offset. This is a minor variation of the PC-relative addressing format.

Addressing Mode Summary


1. Immediate addressing




2. Register addressing




3. Base addressing



4. PC-relative addressing



5. Pseudodirect addressing



C Swap Example

- Illustrates use of assembly instructions for a C bubble sort function

- Swap procedure (leaf)

```
void swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

- v in \$a0, k in \$a1, temp in \$t0

The Procedure Swap

swap:	sll \$t1, \$a1, 2	# \$t1 = k * 4
	add \$t1, \$a0, \$t1	# (address of v[k])
	lw \$t0, 0(\$t1)	# \$t0 (temp) = v[k]
	lw \$t2, 4(\$t1)	# \$t2 = v[k+1]
	sw \$t2, 0(\$t1)	# v[k] = \$t2 (v[k+1])
	sw \$t0, 4(\$t1)	# v[k+1] = \$t0 (temp)
	jr \$ra	# return to calling routine

Assembler Pseudoinstructions

- Most assembler instructions represent machine instructions one-to-one
- Pseudoinstructions:

`move $t0, $t1` \rightarrow `add $t0, $zero, $t1`

`b1t $t0, $t1, L` \rightarrow `slt $at, $t0, $t1`
 `bne $at, $zero, L`

- `$at` (register 1): assembler temporary

Arrays vs. Pointers

- Array indexing involves
 - Multiplying index by element size
 - Adding to array base address
- Pointers correspond directly to memory addresses
 - Can avoid indexing complexity

Example: Clearing and Array

```
clear1(int array[], int size) {
    int i;
    for (i = 0; i < size; i += 1)
        array[i] = 0;
}
```

```
        move $t0,$zero    # i = 0
loop1: sll $t1,$t0,2      # $t1 = i * 4
        add $t2,$a0,$t1   # $t2 =
                        # &array[i]
        sw $zero, 0($t2)  # array[i] = 0
        addi $t0,$t0,1    # i = i + 1
        slt $t3,$t0,$a1   # $t3 =
                        # (i < size)
        bne $t3,$zero,loop1 # if (...)
                        # goto loop1
```

```
clear2(int *array, int size) {
    int *p;
    for (p = &array[0]; p < &array[size];
        p = p + 1)
        *p = 0;
}
```

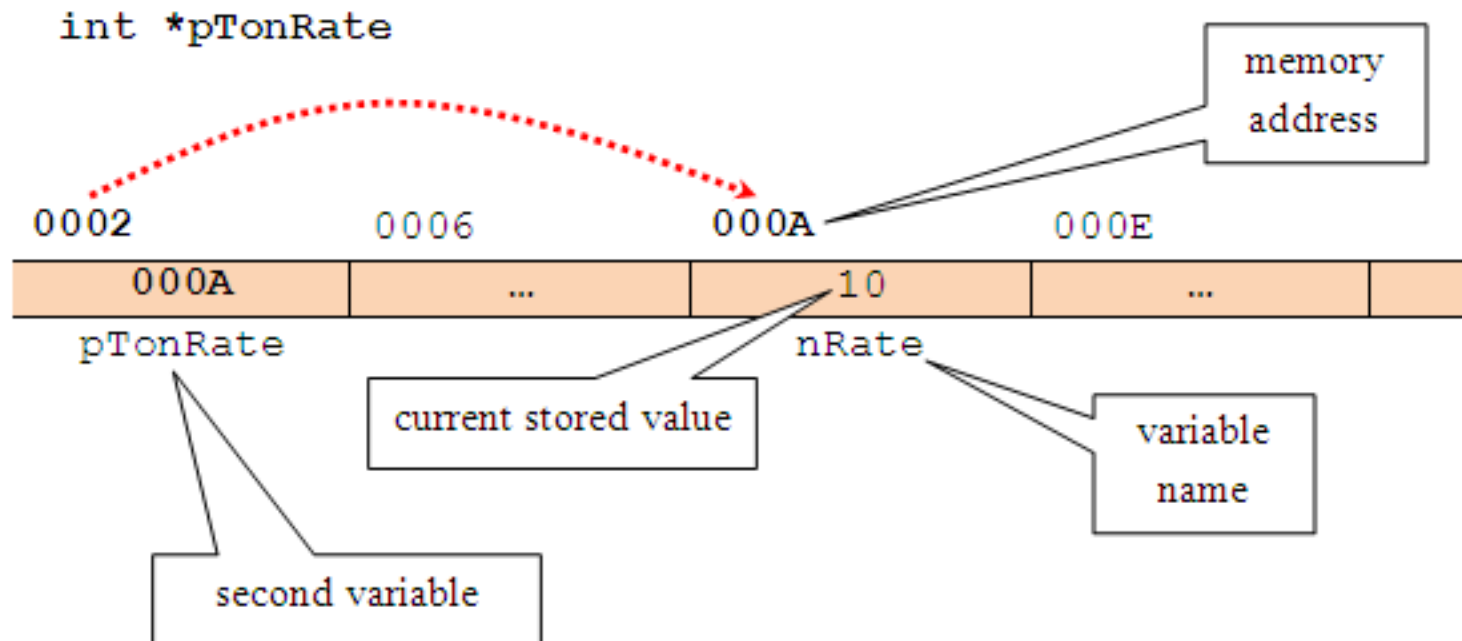
```
        move $t0,$a0      # p = & array[0]
        sll $t1,$a1,2      # $t1 = size * 4
        add $t2,$a0,$t1   # $t2 =
                        # &array[size]
loop2: sw $zero,0($t0)    # Memory[p] = 0
        addi $t0,$t0,4     # p = p + 4
        slt $t3,$t0,$t2   # $t3 =
                        # (p < &array[size])
        bne $t3,$zero,loop2 # if (...)
                        # goto loop2
```

What is a pointer

- In a generic sense, a “pointer” is anything that tells us where something can be found.
- When declaring a variable, the compiler sets aside memory storage with a unique address to store that variable.
- The compiler associates that address with the variable’s name.
- We can manipulate the memory address by using pointers which means that we create a second variable for storing the memory address.

Pointers

- Let us store the `nRate`'s memory address, in `pTonRate` variable.
- So, `pTonRate` now holds the `nRate`'s memory address, where the actual data (10) is stored.
- Pointer variable declaration becomes something like this,
`int *pTonRate;`
- The asterisk (*) is used to show that it is the pointer variable instead of normal variable.



Pointers

- A variable name *directly* references a value.
- A pointer *indirectly* references a value. Referencing a value through a pointer is called *indirection*.
- A pointer variable must be declared before it can be used.
- C uses two pointer operators,
 1. Indirection operator (***) – asterisk symbol, has been explained previously.
 2. Address-of-operator (*&*) – ampersand symbol, means return the address of...
- When *&* operator is placed before the name of a variable, it will return the memory address of the variable instead of stored value.


Pointers

```
int x = 1, y = 2, z[10];  
int *ip;           /* ip is a pointer to an int */  
  
ip = &x;           /* ip points to (contains the memory  
                   address of) x */
```

- **&** "address operator" which gives or produces the memory address of a data variable
- ***** "dereferencing operator" which provides the contents in the memory location specified by a pointer
- The ***** operator is a complement of **&** operator.

Pointers

```
int * pToInt;  
pToInt = &nLocation;  
nLocation = 100;
```



Example: Clearing and Array

```
clear1(int array[], int size) {
    int i;
    for (i = 0; i < size; i += 1)
        array[i] = 0;
}
```

```
        move $t0,$zero    # i = 0
loop1: sll $t1,$t0,2      # $t1 = i * 4
        add $t2,$a0,$t1   # $t2 =
                        # &array[i]
        sw $zero, 0($t2)  # array[i] = 0
        addi $t0,$t0,1    # i = i + 1
        slt $t3,$t0,$a1   # $t3 =
                        # (i < size)
        bne $t3,$zero,loop1 # if (...)
                        # goto loop1
```

```
clear2(int *array, int size) {
    int *p;
    for (p = &array[0]; p < &array[size];
        p = p + 1)
        *p = 0;
}
```

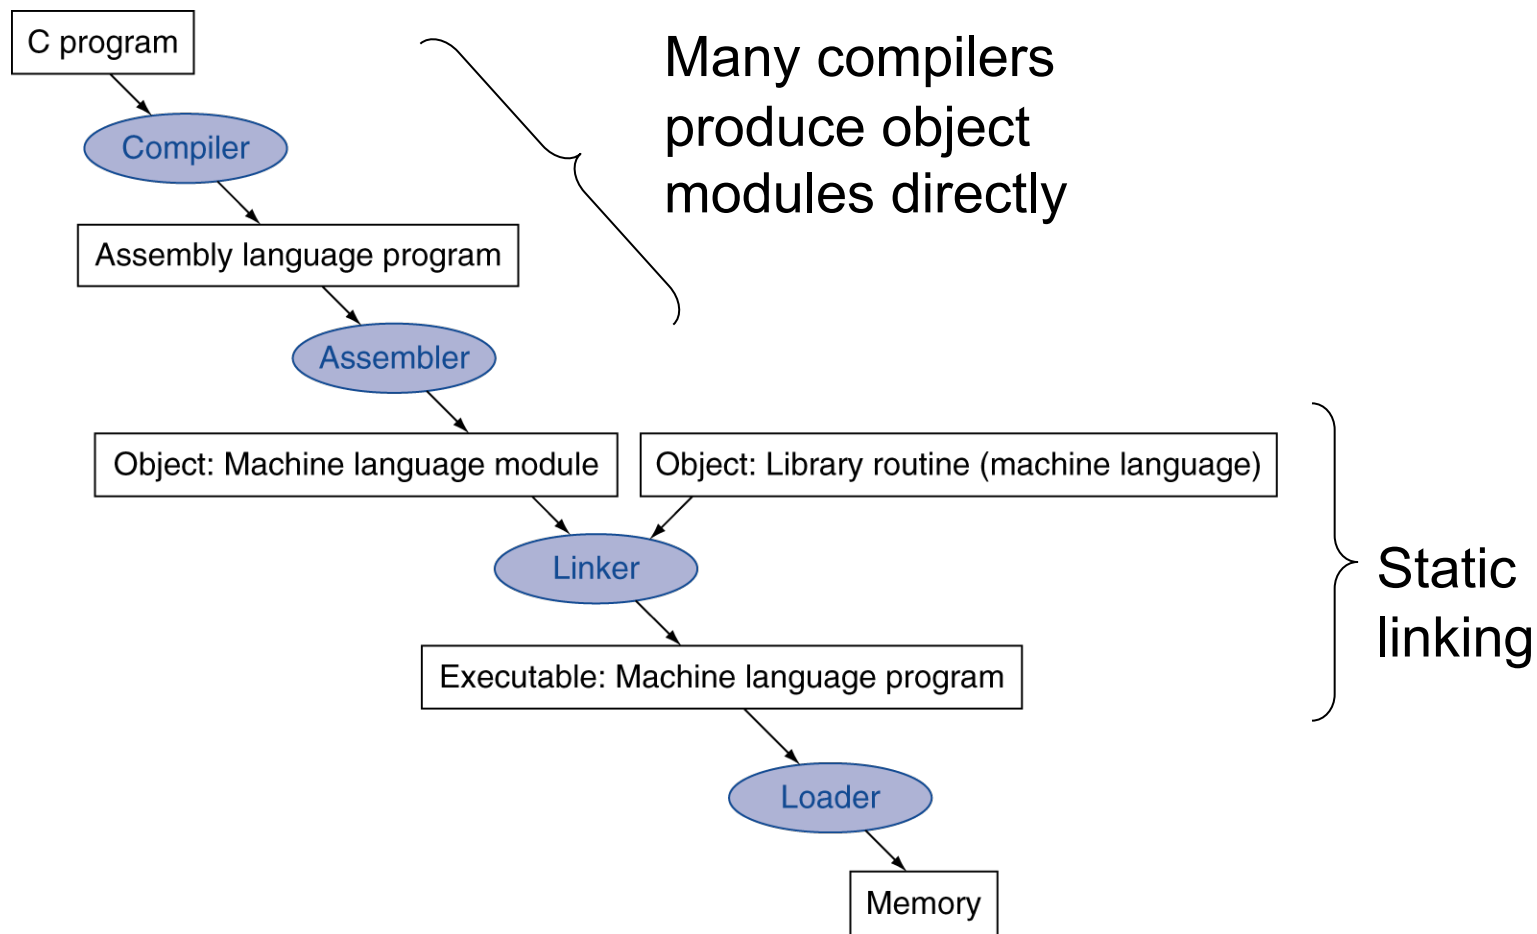
```
        move $t0,$a0      # p = & array[0]
        sll $t1,$a1,2      # $t1 = size * 4
        add $t2,$a0,$t1   # $t2 =
                        # &array[size]
loop2: sw $zero,0($t0)    # Memory[p] = 0
        addi $t0,$t0,4     # p = p + 4
        slt $t3,$t0,$t2   # $t3 =
                        # (p < &array[size])
        bne $t3,$zero,loop2 # if (...)
                        # goto loop2
```

Lecture 10

CMPEN 331

A decorative blue curved shape that starts as a thin line on the left and curves downwards and to the right, ending as a solid blue area in the bottom right corner of the slide.

Translation and Startup



Producing an Object Module

- Assembler (or compiler) translates program into machine instructions
- Provides information for building a complete program from the pieces
 - **Header**: described contents of object module
 - **Text segment**: translated instructions
 - **Static data segment**: data allocated for the life of the program
 - **Relocation info**: for contents that depend on absolute location of loaded program
 - **Symbol table**: global definitions
 - **Debug info**: for associating with source code

Linking Object Modules

- Produces an executable image
 1. Merges segments
 2. Resolve labels (determine their addresses)
 3. Patch location-dependent and external refs

Loading a Program

- Load from image file on disk into memory
 1. Read header to determine segment sizes
 2. Create virtual address space
 3. Copy text and initialized data into memory
 - Or set page table entries so they can be faulted in
 4. Set up arguments on stack
 5. Initialize registers (including \$sp, \$fp, \$gp)
 6. Jump to startup routine
 - Copies arguments to \$a0, ... and calls main
 - When main returns, do exit system call

Dynamic Linking


- Only link/load library procedure when it is called
 - Requires procedure code to be relocatable
 - Avoids image bloat caused by static linking of all (transitively) referenced libraries
 - Automatically picks up new library versions

ARM & MIPS Similarities

- ARM: the most popular embedded core
- Similar basic set of instructions to MIPS

	ARM	MIPS
Date announced	1985	1985
Instruction size	32 bits	32 bits
Address space	32-bit flat	32-bit flat
Data alignment	Aligned	Aligned
Data addressing modes	9	3
Registers	15 × 32-bit	32 × 32-bit
Input/output	Memory mapped	Memory mapped

Instruction Encoding



The Intel x86 ISA

- Evolution with backward compatibility
 - 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
 - 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
 - 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
 - 80286 (1982): 24-bit addresses
 - Segmented memory mapping and protection
 - 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments

The Intel x86 ISA

- Further evolution...
 - i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
 - Pentium (1993): superscalar, 64-bit datapath
 - Later versions added MMX (Multi-Media eXtension) instructions
 - The infamous FDIV bug
 - Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
 - Pentium III (1999)
 - Added SSE (Streaming SIMD Extensions) and associated registers
 - Pentium 4 (2001)
 - New microarchitecture
 - Added SSE2 instructions

The Intel x86 ISA

- And further...
 - **AMD64 (2003): extended architecture to 64 bits**
 - EM64T – Extended Memory 64 Technology (2004)
 - AMD64 adopted by Intel (with refinements)
 - Added SSE3 instructions
 - Intel Core (2006)
 - Added SSE4 instructions, virtual machine support
 - **AMD64 (announced 2007): SSE5 instructions**
 - **Intel declined to follow, instead...**
 - Advanced Vector Extension (announced 2008)
 - Longer SSE registers, more instructions
- If Intel didn't extend with compatibility, its competitors would!

Basic x86 Registers

Name	31	0	Use
EAX			GPR 0
ECX			GPR 1
EDX			GPR 2
EBX			GPR 3
ESP			GPR 4
EBP			GPR 5
ESI			GPR 6
EDI			GPR 7
	CS		Code segment pointer
	SS		Stack segment pointer (top of stack)
	DS		Data segment pointer 0
	ES		Data segment pointer 1
	FS		Data segment pointer 2
	GS		Data segment pointer 3
EIP			Instruction pointer (PC)
EFLAGS			Condition codes

x86 Instruction Encoding

a. JE EIP + displacement

4	4	8
JE	Condition	Displacement

b. CALL

8	32
CALL	Offset

c. MOV EBX, [EDI + 45]

6	1	1	8	8
MOV	d	w	r/m Postbyte	Displacement

d. PUSH ESI

5	3
PUSH	Reg

e. ADD EAX, #6765

4	3	1	32
ADD	Reg	w	Immediate

f. TEST EDX, #42

7	1	8	32
TEST	w	Postbyte	Immediate

- Variable length encoding
 - Postfix bytes specify addressing mode
 - Prefix bytes modify operation
 - Operand length, repetition, locking, ...

x86 Typical Operation

Instruction	Meaning
Control	Conditional and unconditional branches
jnz, jz	Jump if condition to EIP + 8-bit offset; JNE (for JNZ), JE (for JZ) are alternative names
jmp	Unconditional jump—8-bit or 16-bit offset
call	Subroutine call—16-bit offset; return address pushed onto stack
ret	Pops return address from stack and jumps to it
loop	Loop branch—decrement ECX; jump to EIP + 8-bit displacement if ECX ≠ 0
Data transfer	Move data between registers or between register and memory
move	Move between two registers or between register and memory
push, pop	Push source operand on stack; pop operand from stack top to a register
les	Load ES and one of the GPRs from memory
Arithmetic, logical	Arithmetic and logical operations using the data registers and memory
add, sub	Add source to destination; subtract source from destination; register-memory format
cmp	Compare source and destination; register-memory format
shl, shr, rcr	Shift left; shift logical right; rotate right with carry condition code as fill
cbw	Convert byte in eight rightmost bits of EAX to 16-bit word in right of EAX
test	Logical AND of source and destination sets condition codes
inc, dec	Increment destination, decrement destination
or, xor	Logical OR; exclusive OR; register-memory format
String	Move between string operands; length given by a repeat prefix
movs	Copies from string source to destination by incrementing ESI and EDI; may be repeated
lods	Loads a byte, word, or doubleword of a string into the EAX register

Arithmetic for Computers


Chapter 3

Arithmetic for Computers

- Operations on integers
 - Addition and subtraction
 - Multiplication and division
 - Dealing with overflow
- Floating-point real numbers
 - Representation and operations

Integer Addition

- Example: $7 + 6$



- **Overflow if result out of range**
 - Adding +ve and -ve operands, no overflow
 - Adding two +ve operands
 - Overflow if result sign is 1
 - Adding two -ve operands
 - Overflow if result sign is 0

Integer Subtraction

- Add negation of second operand
- Example: $7 - 6 = 7 + (-6)$

$$\begin{array}{r} +7: \quad 0000\ 0000\ \dots\ 0000\ 0111 \\ -6: \quad 1111\ 1111\ \dots\ 1111\ 1010 \\ \hline +1: \quad 0000\ 0000\ \dots\ 0000\ 0001 \end{array}$$

- Overflow if result out of range
 - Subtracting two +ve or two -ve operands, no overflow
 - Subtracting +ve from -ve operand
 - Overflow if result sign is 0
 - Subtracting -ve from +ve operand
 - Overflow if result sign is 1

Dealing with Overflow

- Some languages (e.g., C) ignore overflow
 - Use MIPS `addu`, `addiu`, `subu` instructions
- Other languages (e.g., Ada, Fortran) require raising an exception
 - Use MIPS `add`, `addi`, `sub` instructions
 - On overflow, invoke exception handler
 - Save PC in exception program counter (EPC) register
 - Jump to predefined handler address
 - `mfc0` (move from coprocessor reg) instruction can retrieve EPC value, to return after corrective action